

Growth Machine

Evidence-Driven Outbound Research & Personalization System

WHAT I BUILT

I built an evidence-driven growth machine for a narrow audience: **B2B SaaS documentation and support teams** with public evidence of multilingual, HTML/help-content, template-driven, or other document-heavy workflows.

The machine ingests public company evidence, validates the inputs, detects workflow signals, forms a grounded pain hypothesis, maps only supported SuperDocs capabilities, scores the opportunity, and produces a personalized research/outreach artifact for human review.

The core workflow is: **Public Evidence** → **Validation** → **Workflow Inference** → **Pain Hypothesis** → **Capability Mapping** → **Qualification** → **Personalization** → **Human Review** → **Measurement**.

CALLS I MADE

I deliberately **narrowed the audience** instead of building a generic lead-generation system. The primary channel is outbound research and personalization, but the system **stops before sending outreach**.

I also made **evidence the foundation** of every qualification decision. The system separates what the public evidence actually shows from the workflow and pain hypotheses inferred from it.

WHAT I CUT

I cut automated outreach, private contact enrichment, unrestricted lead scraping, and unsupported product claims. These additions would increase complexity and risk without demonstrating the core thesis: that public evidence can be converted into a specific, reviewable product-fit hypothesis.

HOW QUALIFICATION WORKS

The machine identifies signals such as: HTML/help-center workflows, document editing and revision, repetitive document processes, template-driven document generation, localization and multilingual workflows, and API or programmatic document processing.

These signals are combined into a workflow hypothesis and fit score. Only capabilities supported by the detected evidence are mapped to the account.

MEASUREMENT

The system records discovery count, evidence coverage, qualification rate, fit bands, input fingerprint, and timestamp. Two fresh development runs were executed against different evidence batches. Both produced evidence-backed target records, fit scores, workflow hypotheses, matched capabilities, and personalized artifacts without requiring a code change between runs.

HUMAN LOOP

The output is intentionally a research artifact, not an autonomous sales action. A human must verify: evidence, workflow interpretation, qualification, product-fit claims, personalization, and whether outreach is appropriate. The application cannot send outreach.

HONEST LIMITATIONS

Public evidence is a signal, not inside information. A public feature or workflow does not prove that the company has the hypothesized pain. Workflow inference can be wrong, evidence can become stale, and personalization can still require refinement.

The next meaningful metric would be the percentage of qualified artifacts that a human reviewer considers sufficiently specific and credible to use.

Core Design Principle: *Evidence before inference. Inference before qualification. Qualification before personalization. Human review before action.*